RAG Document Intelligence

Retrieval-Augmented Generation (RAG) combines information retrieval with large language models.

RAG retrieves relevant information from external documents before generating an answer.

This helps language models answer questions using specific document content.

Vector Databases

Vector databases store numerical representations of text called embeddings.

ChromaDB is a vector database that supports semantic similarity search.

Sentence Transformers can generate embeddings locally without requiring a paid API.